

gsDesignNB: Group Sequential Design and Sample Size Re-estimation for Negative Binomial Recurrent Event Trials

Abstract

We present the R package **gsDesignNB** for planning, monitoring, and adapting clinical trials with recurrent event endpoints analyzed using negative binomial models. The package implements sample size calculations that extend the methods of Zhu and Lakkis (2014) and Friede and Schmidli (2010) to handle piecewise accrual, exponential dropout, maximum follow-up truncation, and event gaps. We derive and implement a second-order Taylor correction for Jensen’s inequality bias in the event gap effective rate, reducing the approximation error from 8% to 1% in representative scenarios. The package supports group sequential designs via integration with the **gsDesign** package, and provides both blinded and unblinded sample size re-estimation at interim analyses. A simulation study across 13 scenarios demonstrates that sample size re-estimation effectively recovers power when nuisance parameters are misspecified, provided the treatment effect is as planned, while maintaining Type I error control. The package was developed with substantial assistance from AI coding tools, and we discuss the implications of this workflow for statistical software development.

1 Introduction

Recurrent event endpoints are common in clinical trials for conditions such as multiple sclerosis relapses, COPD exacerbations, and heart failure hospitalizations. The negative binomial (NB) model provides a natural framework for analyzing such data, accounting for both within-subject correlation (overdispersion) and variable follow-up through an exposure offset (Friede and Schmidli 2010; Zhu and Lakkis 2014).

Sample size calculation for NB outcomes requires specifying not only the event rates and treatment effect but also the dispersion parameter, average exposure time, and potentially event gap durations. In practice, several of these parameters may be uncertain at the design stage. Group sequential designs allow for interim monitoring, and sample size re-estimation (SSR) can adapt the trial size when nuisance parameters deviate from assumptions.

The **gsDesignNB** package for R provides an integrated toolkit for these tasks:

1. Sample size and power for fixed designs with flexible enrollment, dropout, follow-up truncation, and event gaps.
2. Group sequential designs via calendar-time analysis schedules compatible with the **gsDesign** package (Jennison and Turnbull 1999).
3. Blinded and unblinded SSR at interim analyses.
4. Simulation utilities for validating designs and evaluating operating characteristics.

This paper describes the methodology, highlights a novel correction for Jensen’s inequality bias in the event gap formula, and presents a simulation study of SSR across multiple scenarios. We also

discuss the role of AI-assisted development in creating the package.

2 Sample size methodology

2.1 Negative binomial model

For subject i in treatment group g with exposure time t_i , the event count Y_i follows a negative binomial distribution with mean $\mu_i = \lambda_g t_i$ and variance $\text{Var}(Y_i) = \mu_i + k\mu_i^2$, where λ_g is the group-specific event rate and k is the dispersion parameter. This arises from a Gamma–Poisson mixture: each subject has a latent rate $\Lambda_i \sim \text{Gamma}(1/k, k\lambda_g)$ and generates events from $\text{Poisson}(\Lambda_i t_i)$.

2.2 Wald test and sample size formula

The treatment effect is the log rate ratio $\theta = \log(\lambda_2/\lambda_1)$, estimated via maximum likelihood under a negative binomial GLM with log link and exposure offset. The Fisher information per subject is $\mathcal{I}_i = \mu_i/(1 + k\mu_i)$, which leads to the sample size formula:

$$n_1 = \frac{(z_{\alpha/s} + z_\beta)^2}{(\theta - \theta_0)^2} \left[\left(\frac{1}{\mu_1} + k_1 \right) + \frac{1}{r} \left(\frac{1}{\mu_2} + k_2 \right) \right]$$

where $\mu_g = \lambda_g \bar{t}_g$, $\bar{t}_g$ is the average exposure, $r = n_2/n_1$, and k_g may be inflated by a variance correction factor $Q_g = E[t_g^2]/(E[t_g])^2$ to account for variable follow-up (Zhu and Lakkis 2014).

This formula is equivalent across the three primary references: Method 3 of Zhu and Lakkis (2014) (fixed sample size), Friede and Schmidli (2010) (blinded SSR), and Mütze et al. (2019) (group sequential designs).

2.3 Jensen’s inequality correction for event gaps

In trials with a mandatory gap of duration γ after each event (e.g., a recovery period), the effective rate is commonly approximated as $\lambda_{\text{eff}} = \lambda/(1 + \lambda\gamma)$. However, this is a population-level average that ignores subject-level heterogeneity. Since $f(x) = x/(1 + x\gamma)$ is concave and subject rates Λ_i are random (Gamma-distributed), Jensen’s inequality gives $E[f(\Lambda)] < f(E[\Lambda])$.

We derive a second-order Taylor correction:

$$\lambda_{\text{eff}} \approx \frac{\lambda}{1 + \lambda\gamma} \left(1 - \frac{k\lambda\gamma}{(1 + \lambda\gamma)^2} \right)$$

using $f''(\lambda) = -2\gamma/(1 + \lambda\gamma)^3$ and $\text{Var}(\Lambda) = k\lambda^2$. Table 1 shows the correction magnitude.

Table 1: Jensen correction magnitude for selected parameter combinations.

λ	k	γ	Correction (%)
1.0	1.0	1.0	25.0
1.0	1.0	0.5	22.2
0.5	1.0	1.0	22.2
0.3	1.0	1.0	17.8

λ	k	γ	Correction (%)
0.5	1.0	0.5	16.0
1.0	0.5	1.0	12.5
0.3	1.0	0.5	11.3
1.0	0.5	0.5	11.1

A simulation study with 10,000 replicates per scenario confirmed that the corrected formula reduces the effective rate estimation error from approximately 8% to 1% and produces slightly conservative power (Table 2).

Table 2: Power comparison: Jensen-corrected vs naive formula (10,000 replicates each).

k	Gap (days)	n (corrected)	n (naive)	Power (corr.)	Power (naive)
0.5	20	436	422	0.9061	0.8992
1.0	20	712	684	0.9170	0.9033
0.5	30	454	436	0.9042	0.8934
1.0	30	740	700	0.9105	0.8988

3 Group sequential design and sample size re-estimation

3.1 Integration with gsDesign

The `gsNBCalendar()` function creates a group sequential design by computing the statistical information at each planned analysis time, accounting for staggered enrollment and variable exposure. It returns an object that inherits from `gsDesign`, enabling use of spending functions, boundary summaries, and integer rounding from the `gsDesign` package.

3.2 Blinded and unblinded SSR

At interim analyses, nuisance parameters (k and λ_1) can be re-estimated to update the required sample size. Blinded SSR (Friede and Schmidli 2010) fits a single NB model to pooled data and splits the estimated rate using the planned treatment effect ratio. Unblinded SSR uses treatment-specific estimates. Both methods accept a `target_info` parameter from the group sequential design to account for inflation due to interim alpha spending.

The `calculate_blinded_info()` function includes a Poisson fallback threshold (matching `mutze_test()`) to guard against numerically unstable dispersion estimates that would produce extreme information values.

3.3 Spending function approach

For the simulation study, bounds are computed using spending functions with `usTime = lsTime = min(planned IF, actual IF)` at the interim. This ensures that neither efficacy nor futility spending exceeds the planned rate, even when more information than expected is available. Futility assessment is deferred until at least 30% of the planned information is observed.

4 Simulation study

4.1 Design

We evaluated SSR across 13 scenarios varying the control rate (λ_1), dispersion (k), accrual rate, and rate ratio ($RR = 0.7, 0.8, 0.9, 1.0$). The planned design targeted 90% power for $RR = 0.7$ with $k = 0.3$, $\lambda_1 = 0.5/\text{month}$, and accrual of 30/month over 12 months. Each non-null scenario used 3,600 simulations; the null scenario used 20,000.

Three strategies were compared: (1) no adaptation, (2) blinded SSR, and (3) unblinded SSR, all using a group sequential design with Hwang–Shih–DeCani spending functions.

4.2 Results

The key findings from the simulation study are:

- **Type I error control:** Under the null ($RR = 1.0$, misspecified nuisance parameters), all three strategies maintained Type I error near or below the nominal 2.5% level (20,000 simulations, $SE \approx 0.1\%$).
- **Power recovery at $RR = 0.7$:** When the planned rate ratio was correct, both SSR methods recovered power from as low as 60% (no adaptation, worst case) to approximately 85–90%, by increasing the sample size up to $2\times$ planned.
- **Diminishing returns at smaller effects:** For $RR = 0.8$ and $RR = 0.9$, the planned sample size (designed for $RR = 0.7$) was inherently insufficient. SSR compensated for nuisance parameter misspecification but could not overcome the smaller treatment effect.
- **Blinded vs unblinded:** Both methods performed similarly across scenarios, supporting the use of blinded SSR to preserve trial integrity.

5 Role of AI in package development

The **gsDesignNB** package was developed with substantial assistance from AI coding tools (Claude, via the Cursor IDE). The AI contributed to:

- **Code generation:** Implementation of simulation engines, data cutting functions, and the group sequential design interface.
- **Mathematical derivation:** The Jensen’s inequality correction was identified and derived through an interactive dialogue analyzing simulation discrepancies. The AI performed Monte Carlo verification and second-order Taylor expansion to pinpoint the bias source.
- **Documentation:** Vignettes, roxygen2 documentation, and notation consistency were developed collaboratively, with the AI drafting and the statistician reviewing for correctness.
- **Debugging:** Several subtle bugs (e.g., `toInteger()` overwriting enrollment with information-scaled values, positional argument mismatches in SSR calls, **gsDesign** collapsing futility bounds when `n.I` is provided directly) were identified through systematic diagnosis.
- **Simulation studies:** The multi-scenario SSR simulation was designed, coded, and iteratively refined through conversation, with the statistician guiding the experimental design and the AI implementing it.

This workflow demonstrates that AI tools can accelerate statistical software development while requiring expert oversight for methodological correctness. The statistician’s role shifts toward

design decisions, mathematical validation, and interpretation, while the AI handles implementation, debugging, and large-scale computation.

6 Discussion

The **gsDesignNB** package fills a gap in the R ecosystem for planning and adapting recurrent event trials with NB outcomes. While the **gscounts** package (Mütze et al. 2019) implements group sequential designs for NB data, **gsDesignNB** extends this with piecewise accrual, event gaps with Jensen’s correction, blinded and unblinded SSR, and comprehensive simulation utilities.

The Jensen’s inequality correction, while modest in many practical scenarios (1–8% effective rate reduction), becomes increasingly important as dispersion and event gap duration grow. The correction is principled (derived from the known Gamma frailty distribution) and conservative (slightly overpowering rather than underpowering).

The simulation study confirms that SSR is an effective tool for maintaining power when nuisance parameters are uncertain, provided the treatment effect assumption is correct. This is consistent with the theoretical result that blinded SSR preserves Type I error while adapting to observed nuisance parameters (Friede and Schmidli 2010).

7 Software availability

The **gsDesignNB** package is available on GitHub at <https://github.com/keaven/gsDesignNB> with documentation at <https://keaven.github.io/gsDesignNB/>. The package requires $R \geq 4.1.0$ and depends on **gsDesign**, **data.table**, **simtrial**, and **MASS**.

References

- Friede, T., and Schmidli, H. (2010), “Blinded sample size reestimation with negative binomial counts in superiority and non-inferiority trials,” *Methods of Information in Medicine*, 49, 618–624. <https://doi.org/10.3414/ME09-02-0060>.
- Jennison, C., and Turnbull, B. W. (1999), *Group sequential methods with applications to clinical trials*, CRC Press.
- Mütze, T., Glimm, E., Schmidli, H., and Friede, T. (2019), “Group sequential designs for negative binomial outcomes,” *Statistical Methods in Medical Research*, 28, 2326–2347. <https://doi.org/10.1177/0962280218773115>.
- Zhu, H., and Lakkis, H. (2014), “Sample size calculation for comparing two negative binomial rates,” *Statistics in Medicine*, 33, 376–387. <https://doi.org/10.1002/sim.5947>.