



Journal of Statistical Software

MMMMMM YYYY, Volume VV, Issue II.

doi: 10.18637/jss.v000.i00

Introduction

Keaven Anderson

Merck & Co., Inc., Rahway, NJ, USA

Hongtao Zhang

Merck & Co., Inc., Rahway, NJ, USA

Andrea Maes

Merck & Co., Inc., Rahway, NJ, USA

Abstract

We present the R package **gsDesignNB** for planning, monitoring, and adapting clinical trials with recurrent event endpoints analyzed using negative binomial models. The package links fixed design sample size calculations, calendar-time group sequential design construction, blinded and unblinded information estimation, sample size re-estimation (SSR), and recurrent-event simulation in a single workflow built on the **gsDesign** package. In the working model, conditional on treatment arm and a subject-level gamma frailty, recurrent events follow a Poisson process with constant within-subject rate, yielding marginal negative binomial counts. For fixed designs, the package extends the methods of [Zhu and Lakkis \(2014\)](#) and [Friede and Schmidli \(2010\)](#) to handle piecewise accrual, piecewise exponential dropout, maximum follow-up truncation, and protocol-defined gaps between counted events. Because statistical information for the log rate ratio depends jointly on exposure, event rates, and dispersion, it is not generally proportional to raw event counts. We derive and implement a second-order delta-method correction for the Jensen-inequality bias in the event-gap effective-rate approximation; in a 64-scenario simulation sweep, the correction reduces representative approximation error from about 8% to 1% and prevents 1–2.5 percentage points of underpowering when dispersion is moderate-to-high and gaps exceed 20 days. A second methodological advance is a score-test workflow with explicit comparison to the traditional Wald/Zhu–Lakkis sample-size formula. The Wald formula uses the alternative variance for both Type I error and power, whereas the score formula uses the null variance for the Type I component and the alternative variance for the power component. A factorial null simulation shows mild finite-sample Type I error inflation for the Wald test and conservative Type I behaviour for the score test; in the displayed superiority scenarios, retaining Wald sizing while using the score test provides a small practical power margin without losing Type I protection. The package also implements a disciplined fallback from maximum likelihood to method-of-moments (MoM) nuisance-parameter estimation: when the NB ML fit is near-Poisson the software switches to a Poisson Wald or score test, but under extreme overdispersion or ML non-convergence it switches to a MoM-based NB test using the observed Fisher information formula, avoiding the anti-conservative behaviour of a blind Poisson fallback. The package also supports calendar-time group sequential monitoring and both blinded and unblinded SSR. A 90-scenario SSR article illustrates the adaptive simulation engine, and production-scale null simulations show that SSR should be checked under the null with the planned final test statistic. In our examples, the Wald test is mildly anti-conservative after SSR, whereas the score test preserves Type I error more conservatively. The paper emphasizes both the statistical methodology and the software architecture that make these workflows reproducible.

Keywords: —!!!—at least one keyword is required—!!!—.

Recurrent event endpoints are common in clinical trials for conditions such as multiple sclerosis relapses, COPD exacerbations, and heart failure hospitalizations. When every patient has the same underlying event rate, a Poisson model suffices: the variance equals the mean, and information accrues linearly with exposure. Real patient populations, however, exhibit substantial between-patient heterogeneity in event rates. Sicker patients contribute many events while most contribute few, inflating the observed variance well beyond the Poisson mean. Ignoring this overdispersion leads to anticonservative tests and overstated power, understating the sample size needed to detect a clinically meaningful treatment effect.

The negative binomial (NB) model addresses this by adding a single dispersion parameter k that captures the excess variation. Published NB fits from pivotal or large trials confirm that $k > 0$ is the rule rather than the exception: estimates commonly fall in the 0.3–1.5 range across disease areas where recurrent events are the primary endpoint, including COPD exacerbations (Wedzicha, Banerji, Chapman, Vestbo, Roche, Ayers, Thach, Fogel, Patalano, and Vogelmeier 2016; Lipson, Barnhart, Brealey, Brooks, Criner, Day, Dransfield, Halpin, Han, Jones *et al.* 2018), heart failure hospitalizations (McMurray, Packer, Desai, Gong, Lefkowitz, Rizkala, Rouleau, Shi, Solomon, Swedberg, and Zile 2014), and multiple sclerosis relapses (Hauser, Bar-Or, Comi, Giovannoni, Hartung, Hemmer, Lublin, Montalban, Rammohan, Selmaj *et al.* 2017). Even moderate dispersion materially reduces the Fisher information for the log rate that each subject contributes — from μ_i under the Poisson to $\mu_i/(1 + k\mu_i)$ under the NB — so that a design ignoring k can understate the required sample size by 10–30%.

Negative binomial design work also becomes operationally more complicated once a trial moves beyond a fixed final analysis. Information for the log rate ratio depends jointly on exposure, event rates, and dispersion, so raw event counts are only an incomplete proxy for monitoring. Accrual may be piecewise, dropout may truncate follow-up, interim looks may be scheduled on the calendar rather than exactly on information fractions, and some protocols only count events separated by a minimum patient-level gap. In software, these trial features must be represented consistently in planning formulas, recurrent-event simulation, interim data cuts, and alpha-spending calculations.

The **gsDesignNB** package for R provides an integrated toolkit for these tasks:

1. Sample size and power for fixed designs with flexible enrollment, dropout, follow-up truncation, protocol event gaps, and either Wald or score-test sizing.
2. Calendar-time group sequential designs that inherit spending-function infrastructure from the **gsDesign** package (Jennison and Turnbull 1999).
3. Blinded and unblinded information estimation together with blinded and unblinded SSR, with a disciplined maximum likelihood / method-of-moments fallback scheme that keeps Wald and score inference well-defined under extreme overdispersion or ML non-convergence.
4. Recurrent-event simulation, including seasonal rates, event- and information-driven interim cutting rules, and operating-characteristic summaries.

The contribution of **gsDesignNB** is therefore both scientific and computational. Scientifically, the package implements NB design formulas for settings where information is not proportional to events, introduces a Jensen-correction for event-gap planning, adds score-test sizing for designs where null calibration is paramount, and replaces the usual blind Poisson fallback with a method-of-moments NB test under genuine overdispersion. Computationally, it

connects planning objects, calendar-time monitoring rules, subject-level simulated trial data, and interim analysis summaries in one workflow. The remainder of the paper first describes that workflow, including the different notions of “gaps” that arise in recurrent-event trials, and then presents the methodological and simulation results that motivated the software design.

Package overview and workflow

1. Scope and relation to gsDesign

`sample_size_nbinom()` is the fixed-design entry point. It returns a `sample_size_nbinom_result` object containing the planning assumptions, expected exposure, expected events, and variance of the log rate ratio. `gsNBCalendar()` lifts that object into a `gsNB` design that inherits from `gsDesign`, so the NB-specific information calculations live in `gsDesignNB` while spending functions, boundary summaries, and integer rounding remain compatible with the underlying `gsDesign` ecosystem. For simulation, `nb_sim()` and `nb_sim_seasonal()` generate subject-level recurrent-event histories, `cut_data_by_date()` and related helpers convert them to interim analysis data sets, and `sim_gs_nbinom()` or `sim_ssr_nbinom()` produce operating characteristics that can be summarized with `summarize_gs_sim()` and `summarize_ssr_sim()`.

Task	Primary functions	Main outputs
Fixed design planning	<code>sample_size_nbinom()</code> , <code>compute_info_at_time()</code>	<code>sample_size_nbinom_result</code> with Wald or score-test sizing
Calendar-time GS design	<code>gsNBCalendar()</code> , <code>toInteger()</code> , <code>gsBoundSummary()</code>	<code>gsNB</code> object with information and calendar timing
Recurrent-event data generation	<code>nb_sim()</code> , <code>nb_sim_seasonal()</code>	Subject-level event histories
Interim data construction	<code>cut_data_by_date()</code> , <code>get_cut_date()</code> , <code>get_analysis_date()</code> , <code>cut_completers()</code>	One row per randomized subject with events and exposure
Interim inference and information	<code>mutze_test()</code> , <code>calculate_blinded_info()</code> , <code>estimate_nb_mom()</code>	Wald or score statistics and blinded/unblinded information estimates
GS operating-characteristic simulation	<code>sim_gs_nbinom()</code> , <code>check_gs_bound()</code> , <code>summarize_gs_sim()</code>	Boundary-crossing summaries
Adaptive simulation	<code>blinded_ssr()</code> , <code>unblinded_ssr()</code> , <code>sim_ssr_nbinom()</code> , <code>summarize_ssr_sim()</code>	Adapted sample sizes and SSR summaries

The package also re-exports the common `gsDesign` spending functions so that users can stay in one namespace when they move from fixed design planning to spending-function

specification. Internally, the simulation engines rely on **data.table** for aggregation and on reproducible parallel random number streams for larger Monte Carlo studies.

2. Worked design example

To make the workflow concrete, consider a semi-real recurrent-event superiority trial inspired by exacerbation studies. The control-arm event rate is 0.5 events/month, the target rate ratio is 0.70, the dispersion is $k = 0.5$, maximum follow-up is 12 months, annual dropout is 10%, and the endpoint counts only events separated by at least 20 days. Accrual is piecewise, with 12 participants/month for the first 6 months and 24 participants/month for the next 6 months. The planned analysis uses the score test with one-sided $\alpha = 0.025$ and 90% power.

```
ex_fixed <- sample_size_nbinom(
  lambda1 = 0.5,
  lambda2 = 0.35,
  dispersion = 0.5,
  power = 0.90,
  alpha = 0.025,
  accrual_rate = c(12, 24),
  accrual_duration = c(6, 6),
  trial_duration = 24,
  dropout_rate = 0.10 / 12,
  max_followup = 12,
  event_gap = 20 / 30.4375,
  test_type = "score"
)

ex_gs <- ex_fixed |>
  gsNBCalendar(
    k = 3,
    test.type = 4,
    alpha = 0.025,
    sfu = sfHSD,
    sfupar = -2,
    sfl = sfCauchy,
    sflpar = c(0.4, 0.75, 0.56, 0.63),
    analysis_times = c(9, 14, 24)
  ) |>
  gsDesignNB::toInteger()
```

Quantity	Value
Fixed-design sample size	268
Group sequential final sample size	324
Target information	99.8
Expected final counted events	1117.8
Average final exposure	11.4

Quantity	Value
Score-test sizing	score

Table 2: Summary of the semi-real recurrent-event design example.

The fixed design requires 268 participants. After calendar-time group sequential monitoring and integer rounding, the final maximum sample size is 324. The same object stores the expected information, exposure, and event totals at each analysis. These quantities are the inputs used by the simulation engines, so the operating-characteristic calculations inherit the same calendar and recurrent-event assumptions.

Look	Month	Information fraction	Cumulative N	Expected events	Upper Z	Lower Z
IA 1	9	0.323	216	239.1	2.691	0.423
IA 2	14	0.726	324	651.6	2.310	0.912
Final	24	1.000	324	1117.8	2.076	2.073

Table 3: Calendar-time group sequential looks for the design example.




Figure 1: Information fraction and cumulative sample size for the design example.

The worked example is intentionally small enough to render as part of the manuscript, but it

is not a toy calculation: it uses piecewise accrual, dropout, maximum follow-up, an event gap, score sizing, calendar-time monitoring, non-binding futility, and integer rounding. The longer vignettes reuse the same objects for simulation and SSR operating-characteristic studies.

3. Timing, information, and three notions of gaps

The package documentation uses the word “gap” in three different senses, and distinguishing them is important for both the methods and the software.

Concept	Scale	Role in the package
Analysis gap Δ_i	Trial calendar	Minimum time between database locks or interim looks
Timing heterogeneity	Accrual, dropout, and follow-up	Determines how exposure and information accumulate over calendar time
Event gap g	Patient time	Defines which recurrent events count and how much post-event time is removed from the at-risk set

The first two affect *when* an analysis can occur. The third affects *which* events contribute to the endpoint and how much at-risk exposure remains after each counted event. In **gs-DesignNB**, the same `event_gap` argument is carried through planning, simulation, and interim data cutting, so the effective event rate, counted events, and at-risk exposure stay aligned. By contrast, calendar-time analysis constraints are handled through `analysis_times` in `gsNBCalendar()` or through the search criteria in `get_cut_date()`.

4. Primary and secondary workflows

For most users, the package workflow is:

1. Use `sample_size_nbinom()` to build a fixed design under the planned rates, dispersion, accrual, dropout, follow-up assumptions, and planned final test statistic.
2. Convert that result to a group sequential design with `gsNBCalendar()`, optionally rounding with `toInteger()`.
3. Choose the appropriate simulator: `sim_gs_nbinom()` for calendar-time GS verification or `sim_ssr_nbinom()` for adaptive simulations with blinded or unblinded SSR.
4. Summarize operating characteristics with `check_gs_bound()`, `summarize_gs_sim()`, and `summarize_ssr_sim()`.

The package also includes several secondary workflows that are important in practice but do not change the core architecture: seasonal recurrent-event simulation, completer-driven interim looks, non-inferiority and super-superiority designs, blinded-information diagnostics, and an interactive Shiny front end to the SSR engine. These are enumerated with their corresponding vignettes in Section 6.3.

5. Dropout, MAR, and imputation scope

Dropout is represented in **gsDesignNB** as exposure truncation. At the planning stage, `sample_size_nbinom()` accounts for piecewise exponential dropout and maximum follow-up through the expected exposure $\bar{t}_g$ in each treatment arm; at interim analyses, data-cutting functions retain each subject's observed counted events and observed time at risk, including partial follow-up before dropout or the analysis cut. Thus the default analysis is not a complete-case analysis that discards subjects who discontinue early. It is an observed-exposure likelihood analysis for recurrent-event counts.

Under an ignorable missing-data process – missing at random (MAR) together with distinct parameters for the observation and outcome processes – direct likelihood inference can use the observed data without modelling the missingness mechanism (Rubin 1976). For the NB log-rate-ratio estimand used here, this means that MAR dropout does not by itself require multiple imputation (MI): the event count accrued before dropout and its corresponding exposure time already contribute to the NB likelihood and to the Wald variance through the offset and information formula. MI may still be useful as an auxiliary reporting or modelling device when important covariates are incompletely observed, but it is not needed simply to make the MAR recurrent-event analysis valid.

Missing not at random (MNAR) dropout is different. If the post-dropout event process depends on unobserved future outcomes or unmeasured deterioration, the observed-exposure likelihood no longer identifies the same full-follow-up estimand without additional assumptions. In that setting MI is best treated as a sensitivity-analysis framework rather than as a default primary analysis (National Research Council 2010). For recurrent events, the controlled-imputation approach of Keene, Roger, Hartley, and Kenward (2014) is a useful template: it uses a negative multinomial formulation, compatible with the Gamma–Poisson view of NB counts, to impute missing post-withdrawal recurrent events under reference-based assumptions such as borrowing the control-arm event behaviour for active-arm dropouts. This targets de facto sensitivity estimands rather than replacing the MAR primary analysis.

A useful future extension would therefore preserve the censoring reason and planned follow-up horizon in the simulated data, then impute post-dropout counts or event times from an NB recurrent-event model with treatment, exposure, prior event history, dropout reason, and prespecified reference-based or delta-adjusted assumptions for informative dropout. Each imputed data set would be analyzed with the same NB Wald estimand, with variance estimation checked as part of the sensitivity analysis. When `event_gap` is part of the estimand, imputing event times rather than only total counts would be preferable so that the protocol counting rule can be re-applied consistently.

Sample size methodology

6. Negative binomial model

For subject i in treatment arm g with exposure time t_i , the event count Y_i follows a negative binomial distribution with mean $\mu_i = \lambda_g t_i$ and variance $\text{Var}(Y_i) = \mu_i + k\mu_i^2$, where λ_g is the arm-specific event rate and k is the dispersion parameter. This arises from a Gamma–Poisson mixture: each subject has a latent rate $\Lambda_i \sim \text{Gamma}(\text{shape} = 1/k, \text{scale} = k\lambda_g)$, so

$E[\Lambda_i] = \lambda_g$ and $\text{Var}(\Lambda_i) = k\lambda_g^2$, and, conditional on Λ_i , the event count is $\text{Poisson}(\Lambda_i t_i)$. In `MASS::glm.nb()` terminology the shape parameter $\theta_{\text{NB}} = 1/k$.

7. Wald and score-test sample size formulas

The treatment effect is the log rate ratio $\theta = \log(\lambda_2/\lambda_1)$, estimated via maximum likelihood under a negative binomial GLM with log link and exposure offset. The design target is statistical information, $\mathcal{J} = 1/\text{Var}(\hat{\theta})$, not the raw event count alone. The Fisher information per subject is $\mathcal{J}_i = \mu_i/(1 + k\mu_i)$. If $W_g = \sum_{i \in g} \mu_i/(1 + k_g\mu_i)$, then

$$\mathcal{J} = \frac{1}{1/W_1 + 1/W_2}.$$

When exposures are summarized by arm-specific averages and the randomization ratio is $r = n_2/n_1$, this leads to the fixed-design sample size formula

$$n_1 = \frac{(z_\alpha + z_\beta)^2}{(\theta - \theta_0)^2} \left[\left(\frac{1}{\mu_1} + k_1 \right) + \frac{1}{r} \left(\frac{1}{\mu_2} + k_2 \right) \right]$$

where $\mu_g = \lambda_g \bar{t}_g$, $\bar{t}_g$ is the average exposure, and k_g may be inflated by a variance correction factor $Q_g = E[t_g^2]/(E[t_g])^2$ to account for variable follow-up (Zhu and Lakkis 2014). Accordingly, two analyses with similar event totals can carry different information if follow-up or dispersion differs.

This is the default `test_type = "wald"` calculation in `sample_size_nbinom()`. It is equivalent across the three primary references: Method 3 of Zhu and Lakkis (2014) (fixed sample size), Friede and Schmidli (2010) (blinded SSR), and Mütze, Glimm, Schmidli, and Friede (2019) (group sequential designs). It is therefore the directly matched planning formula when the final analysis will be the Wald log-rate-ratio test. The practical difference from more familiar first-event formulas is that when the protocol imposes an event gap, the expected count μ_g must reflect the gap-corrected effective rate rather than a naive event-count proxy.

For a score test evaluated under the null, using only the alternative variance for the Type I error component is less natural. `sample_size_nbinom(test_type = "score")` therefore implements a two-variance calculation in the spirit of restricted-null score-test sample sizing (Farrington and Manning 1990):

$$n_1 = \frac{(z_{\alpha/s}\sqrt{V_0} + z_\beta\sqrt{V_1})^2}{(\theta - \theta_0)^2},$$

where

$$V_1 = \left(\frac{1}{\mu_1} + k_1 \right) + \frac{1}{r} \left(\frac{1}{\mu_2} + k_2 \right)$$

is the variance factor under the planned alternative and V_0 is the corresponding null variance factor under the constrained null rate ratio $\exp(\theta_0) = \text{rr0}$. The alpha term is therefore

calibrated on the null information scale, while the beta term remains calibrated on the planned alternative information scale.

This distinction leads to a practical recommendation. The Zhu–Lakkis / Friede–Schmidli / Mutze formula should remain the baseline for a Wald primary analysis and for moderate-information designs where simulation confirms Type I error. When the sample size is modest, dispersion or event gaps are large, the null margin differs from 1, or adaptive/group sequential decisions make Type I calibration central, the more important move is often to use a score statistic for the analysis and operating-characteristic simulations. The two-variance score formula is useful as a test-aligned alternative and diagnostic, but it is not automatically the most powerful choice. For superiority designs like those studied below, the traditional Wald sample size paired with the score test can be a defensible practical default because it preserves Type I error and provides a small sample-size margin. Final designs should therefore compare Wald and score sizing under the planned score test rather than assuming one formula is uniformly preferable.

In `sample_size_nbinom()`, the same information framework is used for superiority, non-inferiority, and super-superiority designs via the `rr0` argument, and the nuisance parameters may be specified either as common scalars or as arm-specific vectors.

8. Jensen’s inequality correction for event gaps

In trials with a mandatory gap of duration γ after each counted event (e.g., a recovery period or a protocol-defined “new-event” window), the effective counted-event rate for a subject with underlying Poisson rate λ is $\lambda_{\text{eff}} = \lambda/(1 + \lambda\gamma)$ (Cox and Lewis 1966). This follows from elementary renewal theory: expected events per unit calendar time equal the reciprocal of the expected inter-event interval $E[T] = 1/\lambda + \gamma$. However, applying this expression at the population rate is a population-level average that ignores subject-level heterogeneity. Since $f(x) = x/(1 + x\gamma)$ is concave and subject rates Λ_i are random (Gamma-distributed), Jensen’s inequality gives $E[f(\Lambda)] < f(E[\Lambda])$.

A second-order delta-method expansion around λ ,

$$E[f(\Lambda)] \approx f(\lambda) + \frac{1}{2}f''(\lambda) \text{Var}(\Lambda),$$

with $f''(\lambda) = -2\gamma/(1 + \lambda\gamma)^3$ and $\text{Var}(\Lambda) = k\lambda^2$, yields

$$\lambda_{\text{eff}} \approx \frac{\lambda}{1 + \lambda\gamma} \left(1 - \frac{k\lambda\gamma}{(1 + \lambda\gamma)^2} \right).$$

Table 5 shows the correction magnitude. When the first-order quantity $k\lambda\gamma/(1 + \lambda\gamma)^2$ approaches 0.25, as in the first row, the second-order expansion is near its limit of accuracy and the full Gamma–Poisson expectation would differ modestly from the displayed value; the correction still moves λ_{eff} in the correct direction and reduces simulation-based approximation error (Table 6).

λ	k	γ	Correction (%)
1.0	1.0	1.0	25.0
1.0	1.0	0.5	22.2
0.5	1.0	1.0	22.2
0.3	1.0	1.0	17.8
0.5	1.0	0.5	16.0
1.0	0.5	1.0	12.5
0.3	1.0	0.5	11.3
1.0	0.5	0.5	11.1

Table 5: Jensen correction magnitude for selected parameter combinations.

A simulation study with 10,000 replicates per scenario confirmed that the corrected formula reduces the effective rate estimation error from approximately 8% to 1% and produces slightly conservative power (Table 6). During package development, this correction was needed for simulation-based power to agree with the analytical planning formula when event gaps were non-negligible.

In the implementation, the same gap rule is used by `nb_sim()` and `cut_data_by_date()`: events that fall inside the post-event gap are not counted, and the corresponding gap time is removed from the at-risk set used for information calculations.

k	Gap (days)	n (corrected)	n (naive)	Power (corr.)	Power (naive)
0.5	20	436	422	0.9061	0.8992
1.0	20	712	684	0.9170	0.9033
0.5	30	454	436	0.9042	0.8934
1.0	30	740	700	0.9105	0.8988

Table 6: Power comparison: Jensen-corrected vs naive formula (10,000 replicates each).

8.1. Broader parameter sweep

The four scenarios in Table 6 fix $\lambda_1 = 0.4$ and vary only k and the gap duration. To characterize *when* the correction matters across a wider range of clinically relevant settings, we conducted a broader simulation sweep over three axes: the control event rate λ_1 , the dispersion k , and the event gap γ .

Parameter grid. The sweep covers $4 \times 4 \times 4 = 64$ scenarios defined by:

Parameter	Values	Clinical rationale
Control rate	0.3, 0.5, 1.0, 2.0 events/month	Low to high recurrent-event burden

Parameter	Values	Clinical rationale
Dispersion	0.1, 0.3, 0.5, 1.0	Near-Poisson to substantial patient heterogeneity
Event gap	10, 20, 30, 45 days	Short to extended counting windows

Table 7: Broad sweep parameter grid for Jensen correction impact study.

The remaining design parameters are held constant across all scenarios: $\lambda_2/\lambda_1 = 0.7$ (30% rate reduction), power = 90%, one-sided $\alpha = 0.025$, dropout rate = 0.1/12/month, maximum follow-up = 12 months, trial duration = 24 months, and piecewise accrual (rate R for 0–6 months, $2R$ for 6–12 months). These are fixed because they affect power symmetrically for both the corrected and naive designs, so they shift both power estimates together without meaningfully changing the *gap* between them. The three chosen axes (λ_1 , k , γ) correspond directly to the three terms in the correction factor:

$$1 - \frac{k\lambda_1\gamma}{(1 + \lambda_1\gamma)^2}$$

Rationale for disease-area calibration. The parameter ranges are informed by published recurrent event trials. Heart failure hospitalization trials such as PARADIGM-HF report control rates of approximately 0.3 events/month with dispersion near 0.5 (McMurray *et al.* 2014). COPD exacerbation trials (e.g., FLAME, IMPACT) report rates around 0.5–0.8/month with $k \in [0.5, 1.0]$ and commonly define a new exacerbation as requiring \$ 28symptom – freedays[@wedzicha2016flame; @lipson2018impact]. Multiplesclerosisrelapsetrialsobserveratesof1–2/yearwithmoderate–to–highoverdispersion[@hauser2017ocrelizumab]. Atoneextreme(_1 = 0.3\$, $k = 0.1$, gap = 10 days), the theoretical correction factor is below 0.1% and the naive formula suffices. At the other ($\lambda_1 = 2.0$, $k = 1.0$, gap = 45 days), it exceeds 10%, and the underpowering from the naive formula becomes practically meaningful.

Simulation protocol. For each scenario, we compute both the Jensen-corrected and naive sample sizes, simulate 3,500 trials at the corrected n , and extract results for the first n_{naive} subjects from the same simulated data. This paired design makes the power comparison very precise (SE $\approx$ 0.5 percentage points at 90% power). The aggregated results — scenario parameters, sample sizes, empirical power, paired power differences, and discordant pair counts — are saved for use in downstream analyses.

Results. Table 8 summarizes selected scenarios from the 64-scenario sweep, ordered by the theoretical correction factor. The corrected design achieves power within simulation error of the 90% target across all 64 scenarios (range: 0.896–0.924), while the naive design falls below 90% in the majority of scenarios where $k \geq 0.5$ and the gap exceeds 20 days.

λ_1	k	Gap (days)	n (corr.)	n (naive)	Δn	Power (corr.)	Power (naive)	Δ Power (pp)
1.0	1.0	30	426	406	20	0.911	0.899	1.1
0.5	1.0	45	490	456	34	0.910	0.890	2.1
1.0	1.0	45	446	420	26	0.911	0.896	1.5
2.0	1.0	30	402	388	14	0.911	0.900	1.0
0.5	1.0	30	466	442	24	0.914	0.893	2.1
0.3	1.0	45	542	502	40	0.919	0.903	1.6
2.0	1.0	45	418	402	16	0.911	0.899	1.2
0.3	1.0	30	516	488	28	0.918	0.904	1.4
0.5	0.5	45	300	284	16	0.908	0.891	1.7
0.3	0.5	45	348	332	16	0.911	0.895	1.7
2.0	0.1	45	96	96	0	0.918	0.918	0.0
0.3	0.1	10	162	162	0	0.900	0.900	0.0

Table 8: Selected scenarios from the broad sweep (3,500 replicates each). Rows are ordered from largest to smallest correction factor; the last two rows illustrate scenarios where the correction has no effect.

Three patterns emerge:

1. **The correction is driven by the $k \times \gamma$ interaction.** Dispersion $k \geq 0.5$ combined with gaps ≥ 30 days consistently produces 1–2 percentage points of underpowering from the naive formula, corresponding to 14–40 additional subjects (3–8% of total n).
2. **High event rates do not amplify the power gap as much as high dispersion.** At $k = 0.3$, even $\lambda_1 = 2.0$ with a 45-day gap shows only \$0.6 pp of underpowering. Conversely, at $k = 1.0$, $\lambda_1 = 0.5$ with a 30-day gap shows 2.1 pp — the largest gap in the sweep. This is because the correction factor $k\lambda\gamma/(1 + \lambda\gamma)^2$ saturates in λ but grows linearly in k .
3. **The correction is negligible when either k or γ is small.** All scenarios with $k = 0.1$ show $\Delta n \leq 4$ and $\Delta \text{Power} < 0.8$ pp. Similarly, all 10-day gap scenarios show $\Delta n \leq 10$ regardless of k .

These findings provide practical guidance: the Jensen correction is most valuable for trials in therapeutic areas with substantial patient heterogeneity ($k \geq 0.5$) and protocol-defined event gaps exceeding 20 days, such as COPD exacerbation studies with 28-day new-event windows or heart failure trials with extended hospitalization-free windows.

Group sequential design, monitoring, and sample size re-estimation

9. Calendar-time design construction

The `gsNBCalendar()` function creates a group sequential design from a `sample_size_nbinom_result`

by evaluating the planned design at user-specified calendar analysis times. For each look it recomputes expected enrollment, exposure, events, and variance under the planned accrual and follow-up structure, then stores the corresponding statistical information in the `n.I` slot inherited from `gsDesign`. The resulting `gsNB` object therefore behaves like a `gsDesign` object for spending functions and boundary summaries while retaining the NB-specific quantities needed for simulation and interpretation.

This calendar-time emphasis matters because recurrent-event trials are often governed by operational rules rather than by exact information fractions. If C_i denotes a calendar floor for look i , Δ_i a minimum gap after the prior look, and $\mathcal{I}_i$ the target information for that look, the analysis time can be expressed schematically as

$$\tau_i = \max \left\{ C_i, \tau_{i-1} + \Delta_i, \inf \left\{ t : \hat{\mathcal{I}}(t) \geq \mathcal{I}_i \right\} \right\}.$$

`gsNBCalendar()` addresses the planning side of this problem through `analysis_times`, while `get_cut_date()` applies the same logic to simulated or observed trial data by combining calendar, event-count, completer, and information targets. This is the package’s bridge between textbook information-time group sequential design and the calendar rules encountered in real monitoring charters.

10. Interim data cuts and information estimates

The function `cut_data_by_date()` converts subject-level recurrent-event histories into an interim analysis data set with one row per randomized subject, containing counted events, time at risk, and total follow-up. The same interface is implemented for both `nb_sim()` and `nb_sim_seasonal()`. `get_analysis_date()` solves the special case of event-driven looks, whereas `cut_date_for_completers()` and `cut_completers()` support analyses defined by the number of subjects who have completed their planned follow-up. More generally, `get_cut_date()` takes the maximum of whichever criteria are supplied, allowing hybrid rules such as “no earlier than month 8 and only after 60 blinded information units.”

Because information depends on exposure and dispersion, the package distinguishes raw event totals from the quantities actually used for monitoring. `calculate_blinded_info()` implements the pooled negative binomial fit of Friede and Schmidli (2010), then splits the pooled rate according to the planned rate ratio. `estimate_nb_mom()` supplies method-of-moments estimates that are useful when maximum likelihood fits become unstable in sparse or highly overdispersed interim data.

11. Wald and score inference with a two-sided ML/MoM fallback rule

For unblinded interim analyses, `mutze_test()` fits the negative binomial log-linear model with an offset for exposure and returns either a Wald statistic for the treatment effect (`test_type = "wald"`) or a score statistic evaluated under the null (`test_type = "score"`). The Wald statistic is the historical analysis model used by Mütze *et al.* (2019) and by the package’s earlier GS simulations. The score statistic fits only the null model, so the numerator and variance are both evaluated on the null information scale; this is the test aligned with the

two-variance sample-size formula above. When the fitted NB shape parameter $\theta_{\text{NB}} = 1/\hat{k}$ is well-behaved, the ML Wald or score test is used directly. When it is not, `mutze_test()` applies one of two *statistically principled* fallbacks rather than a blind Poisson substitution:

- **Near-Poisson regime** ($\hat{k} < 1/\text{poisson_threshold}$, default $\hat{k} < 0.02$): the fitted NB is numerically indistinguishable from a Poisson and the function falls back to the corresponding Poisson Wald or score test. This handles the common case of `MASS::glm.nb()` drifting toward its iteration ceiling when the data truly are Poisson.
- **Extreme-overdispersion or ML-failure regime** ($\hat{k} > \text{mom_threshold}$, default $\hat{k} > 20$; or `glm.nb()` non-convergence): the function re-estimates $(\hat{\lambda}_1, \hat{\lambda}_2, \hat{k})$ by method of moments via `estimate_nb_mom()` and plugs those values into the same observed Fisher information formula that the package uses everywhere else,

$$\hat{\mathcal{J}} = \frac{1}{1/W_1 + 1/W_2}, \quad W_g = \sum_{i \in g} \frac{\hat{\mu}_{g,i}}{1 + \hat{k} \hat{\mu}_{g,i}}, \quad \hat{\mu}_{g,i} = \hat{\lambda}_g t_i.$$

The Wald statistic is then $\hat{\theta}/\sqrt{1/W_1 + 1/W_2}$ with $\hat{\theta} = \log(\hat{\lambda}_2/\hat{\lambda}_1)$; the score statistic uses the corresponding null score and null information.

The methodological point of the second branch is straightforward. Under genuine overdispersion ($k \gg 0$), a Poisson fallback underestimates $\text{Var}(\hat{\theta})$ by a factor of $(1 + k\mu)^{-1}$ per subject, producing anti-conservative tests exactly when the NB model is most needed. MoM does not require ML convergence, always returns a nonnegative $\hat{k}$, and, plugged into the NB information formula, yields a variance estimate with the correct overdispersion structure. The same two-branch rule is used in `calculate_blinded_info()` (pooled blinded estimation) and `unblinded_ssr()` (treatment-specific estimation for SSR), so that every entry point in the package handles fit failure in the same statistically sensible way. Each of these functions now returns a `fallback` field ("ml", "mom", or "poisson") so that simulation engines and diagnostic tables can record which estimator was used at each interim.

The simulation engines also expose blinded and unblinded information columns under both ML and MoM estimation (`info*_ml`, `info*_mom`) so that users can stress-test whether a design's monitoring behaviour is sensitive to the nuisance-parameter estimation method. These alternative columns are diagnostics rather than competing estimands; in well-behaved data they agree closely, while in sparse or unstable interim data the MoM columns remain finite when the ML columns degenerate.

12. Blinded and unblinded SSR

At interim analyses, nuisance parameters (k and the underlying event rates) can be re-estimated to update the required sample size. `blinded_ssr()` uses the pooled model of [Friede and Schmidli \(2010\)](#), preserving masking but typically producing a more conservative adapted sample size. `unblinded_ssr()` instead fits treatment-specific rates and a common dispersion to the interim data, while holding the target treatment effect at its planned value for the final sample size calculation.

Both SSR routines account for dropout in the same exposure-based sense as the fixed-design calculation. The interim nuisance estimates use the observed `tte` supplied by `cut_data_by_date()`, so subjects who drop out before the interim contribute their observed exposure and events. The recalculated total sample size is then obtained by calling `sample_size_nbinom()` with the planned accrual, dropout, maximum follow-up, and event-gap assumptions. The resulting SSR therefore adjusts for the loss of information expected from dropout under the specified ignorable censoring model; it is not an MNAR sensitivity procedure.

`sim_ssr_nbinom()` wraps these choices into a full adaptive simulation engine. It supports “No adaptation”, “Blinded SSR”, and “Unblinded SSR” strategies; searches for interim looks using information-based timing; and allows efficacy and futility bounds to be driven by alternative information measures through `bound_info = "unblinded_ml", "blinded_ml", "unblinded_mom", or "blinded_mom"`. In the simulation study below, bounds are computed with spending functions using `usTime = lsTime = min(planned IF, actual IF)` at the interim so that spending is not accelerated when more information than expected is observed. Futility assessment is deferred until at least 30% of the planned information is available.

From a computing perspective, both `sim_gs_nbinom()` and `sim_ssr_nbinom()` are designed for large Monte Carlo studies: they use `data.table` for repeated aggregation, support reproducible parallel execution, and return both trial-level and analysis-level summaries so that design calibration can be checked after the fact.

Reproducible workflows

13. Example 1: fixed design to calendar-time GS simulation

The first workflow shows the main path from fixed-design planning to a calendar-time GS simulation. The full executable version appears in the group-sequential simulation vignette; the code here is a minimal example.

```
library(gsDesignNB)

enroll_rate <- data.frame(rate = 12, duration = 6)
fail_rate <- data.frame(
  treatment = c("Control", "Experimental"),
  rate = c(0.6, 0.42),
  dispersion = c(0.4, 0.4)
)
dropout_rate <- data.frame(
  treatment = c("Control", "Experimental"),
  rate = c(0.05, 0.05),
  duration = c(12, 12)
)

fixed_design <- sample_size_nbinom(
  lambda1 = 0.6,
  lambda2 = 0.42,
```



```

    dispersion = 0.4,
    power = 0.8,
    alpha = 0.025,
    accrual_rate = enroll_rate$rate,
    accrual_duration = enroll_rate$duration,
    trial_duration = 12,
    max_followup = 6
  )

gs_design <- gsNBCalendar(
  fixed_design,
  k = 3,
  test.type = 4,
  analysis_times = c(4, 8, 12)
)

# n_sims = 50 is for illustration only; the GS verification vignette uses
# 3,600 replicates for its cached operating-characteristic results.
sim_gs <- sim_gs_nbinom(
  n_sims = 50,
  enroll_rate = enroll_rate,
  fail_rate = fail_rate,
  dropout_rate = dropout_rate,
  max_followup = 6,
  n_target = ceiling(utils::tail(gs_design$n_total, 1)),
  design = gs_design,
  cuts = lapply(gs_design$T, function(t) list(planned_calendar = t)),
  seed = 20260413
)

oc <- summarize_gs_sim(
  check_gs_bound(sim_gs, gs_design, info_col = "info_unblinded_ml")
)

oc$analysis_summary

```

The `analysis_summary` element of `oc` contains per-look cumulative rejection and futility probabilities, mean observed information, and mean calendar time at each look, so that the simulated operating characteristics can be compared directly to the planned boundaries. This example illustrates the main software-design principle of the package: the fixed design determines the target information, `gsNBCalendar()` maps that information to calendar looks, `sim_gs_nbinom()` generates trial histories consistent with the planning assumptions, and `check_gs_bound()` applies the group sequential boundaries on the selected information scale.

14. Example 2: information-based SSR

The second workflow shows how the same design object can be used in an adaptive simulation with blinded and unblinded sample size re-estimation.

```
library(gsDesignNB)

ssr_res <- sim_ssr_nbinom(
  n_sims = 50,
  enroll_rate = enroll_rate,
  fail_rate = fail_rate,
  dropout_rate = dropout_rate,
  max_followup = 6,
  design = gs_design,
  strategies = c("No adaptation", "Blinded SSR", "Unblinded SSR"),
  bound_info = "unblinded_ml",
  seed = 20260413
)

summarize_ssr_sim(ssr_res)$trial_summary[
  , c("strategy", "rejection_rate", "mean_n", "pct_adapted")
]
```

The SSR simulation-study article uses this same engine to define a 90-scenario grid and adds design constraints on when adaptation is allowed. The bundled package cache keeps the non-null power grid deliberately small for portable documentation builds, while the null Type I calibration tables are stored at production scale. The essential software idea is already visible in the compact example: one **gsNB** planning object can drive both non-adaptive and adaptive operating-characteristic evaluations.

15. Additional workflows

The package includes further worked examples that round out the software contribution.

Workflow	Main functions	What it adds
Seasonal recurrent-event simulation	<code>nb_sim_seasonal()</code> , <code>cut_data_by_date()</code>	Calendar-time variation in failure rates
Completer-based monitoring	<code>cut_date_for_completers()</code> , <code>cut_completers()</code>	Interim looks keyed to completed follow-up
Non-inferiority and super-superiority	<code>sample_size_nbinom(rr0 = ...)</code>	Alternative null hypotheses in the same design framework
Blinded-information diagnostics	<code>calculate_blinded_info()</code> , <code>estimate_nb_mom()</code>	Stability checks for sparse or unstable interims
Verification studies	<code>nb_sim()</code> , <code>mutze_test()</code> , <code>sim_gs_nbinom()</code>	Simulation-based agreement checks for design formulas

These workflows are documented in the package vignettes so that the manuscript examples

can remain brief while the installed package still serves as the full replication resource.

Simulation study

16. Design

The manuscript reports selected operating characteristics, while the package articles serve as executable supplementary material. This keeps the paper readable in a software-journal format but still makes the simulation design, code, and fuller numerical output available to readers. In particular, the `verification-simulation`, `simulation-example`, `score-vs-wald-simulation`, and `ssr-simulation-study` articles contain the code used to reproduce the verification checks, event-gap studies, score-versus-Wald calibration, and SSR operating characteristics summarized below.

The `ssr-simulation-study` article defines a 90-scenario grid varying the control rate $\lambda_1 \in \{0.3, 0.5, 0.8\}$ /month, dispersion $k \in \{0.5, 1.0\}$, planned accrual rate $\in \{12, 18, 24\}$ /month, and true rate ratio $RR \in \{0.6, 0.7, 0.85, 1.0, 1.1\}$ (ninety combinations in total). The planned design was built with `sample_size_nbinom()` and `gsNBCalendar()`, targeting 90% power for $RR = 0.7$ with $k = 0.5$, $\lambda_1 = 0.5$ /month, and planned accrual of 26/month, for a planned total sample size of 312 and a maximum follow-up of 12 months. The bundled article cache uses 3,600 replicates per $RR < 1$ power scenario, 20,000 replicates per $RR = 1$ null scenario, and 1,000 replicates per $RR > 1$ interior-null scenario. The dedicated non-binding Type I tables use an additional 20,000 replicates per dispersion and test statistic. See the `ssr-simulation-study` vignette for the full scenario grid, replicate controls, and cache-regeneration code.

Three strategies were compared: (1) no adaptation, (2) blinded SSR, and (3) unblinded SSR, all using a group sequential design with Hwang–Shih–DeCani spending functions (Hwang, Shih, and de Cani 1990).

Table 10 summarizes how the Monte Carlo evidence is divided between the manuscript and the package supplement. For a rejection probability $\hat{p}$ based on B replicates, the Monte Carlo standard error is $\sqrt{\hat{p}(1 - \hat{p})/B}$; this is reported or bounded for the main simulation summaries.

Component	Manuscript role	Supplementary package article	Replicates and reported quantities
Fixed-design verification	Demonstrate agreement between NB planning formulas and simulated power	<code>verification-simulation</code> , <code>sample-size-nbinom</code>	Scenario-specific power, information, and Monte Carlo error checks

Component	Manuscript role	Supplementary package article	Replicates and reported quantities
Event-gap correction	Show where the Jensen correction changes sample size and power	<code>simulation-example</code> , <code>verification-simulation</code>	64-scenario sweep; 3,500 or 10,000 replicates depending on table; corrected vs naive n and power
Wald versus score inference	Compare Wald and score tests under Wald- and score-sized designs	<code>score-vs-wald-simulation</code>	27-scenario factorial grid; 20,000 null replicates and 3,500 power replicates per cell; Type I error, power, and null Z distributions
SSR Type I calibration	Summarize null rejection and adapted sample size	<code>ssr-simulation-study</code>	20,000 replicates for $RR = 1$; rejection rate, MCSE, mean n , and percent adapted
SSR power recovery	Compare no adaptation, blinded SSR, and unblinded SSR under nuisance misspecification	<code>ssr-simulation-study</code>	Bundled production cache uses 3,600 replicates per $RR < 1$ scenario and 1,000 replicates per $RR > 1$ scenario
SSR starting-size sensitivity	Compare Wald- and score-sized starting designs under the score final test	<code>ssr-simulation-study</code> ; <code>data-raw/generate_ssr_size</code>	Low-event stress setting; 5,000 null replicates and 3,000 power replicates per starting-size rule
Estimation diagnostics	Check whether ML, MoM, and fallback labels affect monitoring conclusions	<code>blinded-info-diagnostics</code> , <code>ssr-simulation-study</code>	Blinded/unblinded information paths, nuisance estimates, and estimator fallback labels

Table 10: Simulation reporting map. The package articles are treated as executable supplementary material for the condensed results reported in the manuscript.

17. Results

Wald versus score calibration. The `score-vs-wald-simulation` article evaluates 27 base scenarios varying $\lambda_1 \in \{0.15, 0.40, 1.00\}$, $k \in \{0.2, 0.5, 1.0\}$, and event gaps of 0, 15, or 30 days, with both Wald- and score-sized designs. In this grid, score sizing changed the total sample size by at most four subjects and was never larger than Wald sizing, because the null variance was smaller than the alternative variance for these superiority assumptions. The empirical null results therefore isolate the effect of the test statistic more than the effect of sample size: Wald-test Type I error ranged from 0.0243 to 0.0317, with 13–15 of 27 scenarios above nominal beyond Monte Carlo error depending on the sizing method; score-test Type I error ranged from 0.0200 to 0.0264, with no scenario above nominal beyond Monte Carlo error and several conservative cells. Under $RR = 0.70$, Wald-test power averaged about 0.904–0.907. Score-test power averaged 0.895 under Wald sizing and 0.893 under score sizing; 8–9 of 27 score-test scenarios were below the 90% target beyond Monte Carlo error. These results support using the score test when Type I error preservation is the priority, while treating Wald sizing as a useful practical baseline or power margin for score-test superiority designs. Score-test designs should still be simulated for power and may need a modest information margin.

SSR Type I calibration. Table 11 summarizes the empirical one-sided Type I error under $RR = 1.0$ for the two null dispersion configurations at nominal $\alpha = 0.025$, comparing the Wald statistic with the score statistic under the same group sequential design. With 20,000 replicates per cell the Monte Carlo standard error is about 0.001. The Wald statistic is mildly anti-conservative, with empirical Type I error ranging from 0.0279 to 0.0312. The score statistic is conservative in the same design, with empirical Type I error ranging from 0.0214 to 0.0232. Thus the preferred remedy is to align the analysis and simulation engine with the score test and then verify power, rather than to lower alpha for a Wald analysis.

Test	k	Strategy	Type I	Mean n	% adapted
Wald	0.5	No adaptation	0.0306	312.0	0.0
Wald	0.5	Blinded SSR	0.0309	320.5	33.2
Wald	0.5	Unblinded SSR	0.0309	325.4	45.3
Wald	1.0	No adaptation	0.0279	312.0	0.0
Wald	1.0	Blinded SSR	0.0307	508.4	98.4
Wald	1.0	Unblinded SSR	0.0312	519.5	98.4
Score	0.5	No adaptation	0.0229	312.0	0.0
Score	0.5	Blinded SSR	0.0229	320.6	33.4
Score	0.5	Unblinded SSR	0.0232	325.4	45.5
Score	1.0	No adaptation	0.0214	312.0	0.0
Score	1.0	Blinded SSR	0.0229	509.5	99.0
Score	1.0	Unblinded SSR	0.0232	520.5	99.0

Table 11: Empirical one-sided Type I error and average adapted sample size under the null ($RR = 1.0$), 20,000 replicates per cell. Futility stopping is ignored in this non-binding calibration table so all trials continue unless they cross an efficacy boundary.

Power recovery at the designed effect ($RR = 0.7$). In the production score-test cache, no-adaptation power ranged from 0.698 to 0.896 across nuisance scenarios and averaged 0.807.

Blinded and unblinded SSR both ranged from about 0.880 to 0.903 and averaged 0.893. The recovery was most important under the higher-dispersion scenarios: for $k = 1.0$, mean no-adaptation power was 0.733, whereas blinded and unblinded SSR averaged 0.889. Because the score test is conservative, designs using it may still need a modest information margin to achieve a 90% power target across the full nuisance range.

Diminishing returns at smaller effects. For $RR = 0.85$ (true effect smaller than planned), the planned design is inherently underpowered; SSR compensates for nuisance-parameter misspecification but cannot overcome the effect-size mismatch. In the production score-test cache, no-adaptation power averaged 0.242, while blinded and unblinded SSR averaged about 0.296 and 0.297. For $RR = 1.1$ (treatment worse than control), all three strategies had maximum rejection rates of 0.005 across the scenario grid, confirming that SSR does not artificially inflate one-sided rejections in the wrong direction.

Blinded versus unblinded SSR. In this score-test configuration, blinded and unblinded SSR have very similar operating characteristics. At $RR = 0.7$ their average power differs by less than 0.001 and their average sample sizes are both about 351 subjects. Under the non-binding null calibration, both strategies adapt frequently when $k = 1.0$, and unblinded SSR has slightly larger mean sample size in the displayed scenarios. The practical conclusion is therefore not that one approach is uniformly more efficient, but that the masked and unmasked update rules should be compared under the trial-specific nuisance range, adaptation cap, and analysis statistic.

SSR starting-size sensitivity. The fixed-design score-versus-Wald sweep suggests that the traditional Wald sample size can provide a small information margin when the planned analysis uses the score test. To check whether that margin carries into SSR, the SSR article adds a targeted low-event stress setting with $\lambda_1 = 0.15$, $k = 0.5$, $RR = 0.7$, no event gap, and the same score final test. The Wald-sized group sequential design has fixed-design $n = 390$ and final group sequential $n = 472$, compared with 384 and 464 for the score-sized design. With 5,000 non-binding null replicates per starting-size rule, score-test Type I error remains close to nominal for both starts: 0.0252–0.0258 for the Wald-sized start and 0.0258–0.0266 for the score-sized start (MCSE about 0.0022). With 3,000 planned-effect replicates per starting-size rule, power is also similar: 0.898–0.907 for the Wald-sized start and 0.900–0.911 for the score-sized start (MCSE about 0.005). Thus Wald sizing should be viewed as a candidate margin to evaluate, not as a guaranteed SSR power improvement.

Discussion

The **gsDesignNB** package fills a gap in the R ecosystem for planning and adapting recurrent-event trials with NB outcomes. Its statistical contributions are the alignment of fixed-design planning, calendar-time monitoring, and sample size re-estimation around Fisher information for the log rate ratio; the Jensen correction for event gaps with accompanying simulation evidence; and a two-sided ML/MoM fallback rule that keeps Wald inference well-defined under both near-Poisson and extreme-overdispersion regimes. Its computing contributions are the object flow from `sample_size_nbinom_result` to `gsNB` to simulation outputs; subject-level recurrent-event simulation with event-gap-aware interim cuts; diagnostic information calculations under ML and method-of-moments nuisance estimation; and reproducible Monte Carlo engines for both fixed and adaptive monitoring strategies.

The Jensen correction is important precisely because the software is designed to keep planning

and simulation on the same scale. When event gaps are non-negligible, the naive renewal approximation overstates the effective rate and therefore overstates information. The 64-scenario sweep confirms that the correction is most valuable when dispersion and event-gap duration are both substantial: at $k \geq 0.5$ with gaps ≥ 30 days, the naive formula underpowers by about 1–2 percentage points, whereas at $k = 0.1$ or gaps ≤ 10 days the correction adds little to the required sample size. Because the same gap rule is carried into simulation and interim cutting, users can see directly when the correction matters for operating characteristics rather than only for algebra.

The score-versus-Wald sweep clarifies how the package should be used for sample-size planning. The Zhu–Lakkis / Friede–Schmidli / Mutze calculation remains the appropriate first-line calculation for a Wald primary analysis, because its variance factor is exactly the Wald alternative variance used by that analysis. It should not, however, be presented as a universal guarantee of finite-sample Type I error control. When the design is lower-information, highly overdispersed, affected by event gaps, based on a non-unity null margin, or adaptively monitored, a score-test primary analysis is often the more defensible default. The corresponding best practice is to simulate the score statistic directly and compare candidate sample-size rules under that statistic. The two-variance score formula uses the null variance V_0 for the alpha component and the planned alternative variance V_1 for the beta component, but in the superiority grid it was slightly smaller than Wald sizing and slightly less powerful for the score test. Thus Wald/Zhu–Lakkis sizing paired with the score test is a reasonable practical recommendation when it supplies a modest information margin. Because the score test can be conservative, the final design should still be calibrated for both Type I error and power; increasing the target power or total information slightly may be preferable to reverting to an anti-conservative Wald test.

The SSR study likewise shows how the package’s computational layer supports statistical design checks. The production score-test grid confirms that adverse nuisance-parameter misspecification can be diagnosed and that SSR moves sample size in the expected direction, but it also shows why the final test statistic matters. The non-binding null simulations show that Wald analyses are mildly anti-conservative in this configuration, whereas score analyses are conservative at the same nominal one-sided alpha. The updated SSR article therefore compares Wald and score analyses directly rather than treating an ad hoc alpha reduction as the primary remedy. The larger adapted sample sizes under high dispersion should be interpreted primarily as nuisance-parameter misspecification and adaptation-cap behavior, not as a benefit of the Wald test. The score test is important because it controls rejection after SSR has increased information; it does not remove the need to simulate how often the adaptation rule increases sample size. The targeted starting-size sensitivity reinforces that distinction: an eight-subject Wald starting-size margin did not yield a clear SSR power gain, so starting-size margins should be evaluated inside the planned adaptation rule. Blinded and unblinded nuisance updates had similar power and sample-size behavior in the score-test configuration, so their relative efficiency should be calibrated for the trial-specific adaptation rule. The ability to compare `bound_info` choices, ML versus method-of-moments information paths, and the new ML/MoM fallback behaviour inside the same simulation engine is a practical advantage of the package design: because each interim reports the fallback label actually used, simulation summaries can verify that adaptive decisions were driven by well-posed estimators rather than degenerate ML fits.

The package’s current missing-data position is deliberately narrower than a full causal estimand framework. Dropout and administrative censoring are handled as exposure truncation under ignorable censoring, which is the setting in which observed-data NB likelihood inference is appropriate and MI is not required. MNAR sensitivity analyses would need additional data structure – in particular, explicit censoring reason, planned remaining follow-up, and assumptions about post-dropout event rates – before MI can be implemented transparently. This is a natural extension of the simulator and data-cutting layer, but it should not be conflated with the primary MAR analysis.

From a software perspective, **gsDesignNB** is complementary to both **gsDesign** and **gscounts**. The general-purpose **gsDesign** package provides mature spending-function and boundary calculations but does not itself supply NB recurrent-event planning formulas, calendar-time translation of those formulas, or recurrent-event simulation with blinded information and SSR. The **gscounts** implementation motivated by Mütze *et al.* (2019) addresses group sequential NB testing, whereas **gsDesignNB** emphasizes the broader workflow needed in practice: piecewise accrual and dropout, maximum follow-up, protocol event gaps with Jensen correction, seasonal simulation, completer- or information-driven interim timing, blinded as well as unblinded SSR simulation, and a two-sided ML/MoM fallback rule for unstable NB ML fits.

The package’s secondary workflows matter for a software paper even when they are not the main empirical case study. Seasonal simulation extends the constant-rate gamma-frailty model to a common time-varying setting; completer-based cuts capture another operational monitoring rule; and the non-inferiority, blinded-diagnostics, and verification vignettes show that the same core objects support multiple design questions. In that sense, the package is not only a collection of isolated functions but a coherent interface for recurrent-event trial planning and evaluation.

A two-sided ML/MoM fallback rule. A further methodological feature delivered in the current release replaces the previously permissive single-sided Poisson fallback with a two-sided rule. Near-Poisson data ($\hat{k} < 0.02$) still fall back to the corresponding Poisson Wald or score test, where the Poisson variance is statistically correct. Under extreme overdispersion ($\hat{k} > 20$) or ML non-convergence, however, the package now falls back to a method-of-moments re-estimation of $(\hat{\lambda}_1, \hat{\lambda}_2, \hat{k})$ via `estimate_nb_mom()` and plugs those values into the observed Fisher information formula that the package uses everywhere else, rather than a Poisson test. This matters because under genuine overdispersion the Poisson variance is anti-conservative by a factor of $(1 + k\mu)^{-1}$ per subject, which is exactly the regime where a robust fallback is most needed. The MoM path requires no ML convergence, always returns a nonnegative $\hat{k}$, and preserves the NB variance structure. The rule is applied uniformly in `mutze_test()`, `calculate_blinded_info()`, and `unblinded_ssr()`, each of which now reports a `fallback` label ("ml", "mom", or "poisson") in its return value so that simulation engines can audit which estimator was used at each interim. In `mutze_test()` the behaviour is controlled by two thresholds, `poisson_threshold` (default 50, i.e., $\hat{k} < 0.02$) and `mom_threshold` (default 20, i.e., $\hat{k} > 20$); both are user-tunable for sensitivity analyses.

Reproducibility and software availability

The **gsDesignNB** package is available on GitHub at <https://github.com/keaven/gsDesignNB> with documentation at <https://keaven.github.io/gsDesignNB/>. The pack-

age requires $R \geq 4.1.0$ and depends primarily on R (R Core Team 2025), **gsDesignNB** (Anderson, Zhang, and Maes 2026), **gsDesign** (Anderson 2026), **data.table** (Barrett, Dowle, Srinivasan, Gorecki, Chirico, Hocking, Schwendinger, and Krylov 2026), **simtrial** (Anderson, Zhao, Blischak, and Zhang 2025), and **MASS** (Venables and Ripley 2002). Figures and rendered package documentation use **ggplot2** (Wickham 2016), **scales** (Wickham, Pedersen, and Seidel 2025b), **knitr** (Xie 2025), **rmarkdown** (Allaire, Xie, Dervieux, McPherson, Luraschi, Ushey, Atkins, Wickham, Cheng, Chang, and Iannone 2025), and **pkgdown** (Wickham, Hesselberth, Salmon, Roy, and Brueggemann 2025a).

The source distribution includes vignettes for fixed-design NB sample size calculation, calendar-time GS simulation, end-to-end SSR, the larger SSR simulation study, blinded-information diagnostics, seasonal simulation, completer-based interims, non-inferiority design, and verification of the planning formulas against simulation. We treat these articles as supplementary material for the paper: the manuscript gives the design rationale, selected tables, and interpretation, while the package articles provide the fuller grids, code paths, and rendered results. The two code examples in this paper are deliberately small extracts of those vignettes so that the manuscript highlights the core software interface while the installed package retains more extensive executable workflows. For local exploratory use, the package also provides `run_ssr_shiny()` as an interactive front end to the SSR engine and `preview_pkgdown_site()` for documentation preview.

References

References

- Allaire J, Xie Y, Dervieux C, McPherson J, Luraschi J, Ushey K, Atkins A, Wickham H, Cheng J, Chang W, Iannone R (2025). *rmarkdown: Dynamic Documents for R*. R package version 2.30, URL <https://github.com/rstudio/rmarkdown>.
- Anderson K (2026). *gsDesign: Group Sequential Design*. R package version 3.9.0.9003, URL <https://keaven.github.io/gsDesign/>.
- Anderson K, Zhang H, Maes A (2026). *gsDesignNB: Sample Size and Simulation for Negative Binomial Outcomes*. R package version 0.3.2, URL <https://keaven.github.io/gsDesignNB/>.
- Anderson K, Zhao Y, Blischak J, Zhang Y (2025). *simtrial: Clinical Trial Simulation*. doi:10.32614/CRAN.package.simtrial. R package version 1.0.2, URL <https://CRAN.R-project.org/package=simtrial>.
- Barrett T, Dowle M, Srinivasan A, Gorecki J, Chirico M, Hocking T, Schwendinger B, Krylov I (2026). *data.table: Extension of data.frame*. doi:10.32614/CRAN.package.data.table. R package version 1.18.2.1, URL <https://CRAN.R-project.org/package=data.table>.
- Cox DR, Lewis PAW (1966). *The Statistical Analysis of Series of Events*. Methuen, London.
- Farrington CP, Manning G (1990). “Test statistics and sample size formulae for comparative binomial trials with null hypothesis of non-zero risk difference or non-unity relative risk.” *Statistics in Medicine*, **9**(12), 1447–1454. doi:10.1002/sim.4780091208.

- Friede T, Schmidli H (2010). “Blinded sample size reestimation with negative binomial counts in superiority and non-inferiority trials.” *Methods of Information in Medicine*, **49**(06), 618–624. doi:[10.3414/ME09-02-0060](https://doi.org/10.3414/ME09-02-0060).
- Hauser SL, Bar-Or A, Comi G, Giovannoni G, Hartung HP, Hemmer B, Lublin F, Montalban X, Rammohan KW, Selmaj K, *et al.* (2017). “Ocrelizumab versus interferon beta-1a in relapsing multiple sclerosis.” *New England Journal of Medicine*, **376**(3), 221–234. doi:[10.1056/NEJMoa1601277](https://doi.org/10.1056/NEJMoa1601277).
- Hwang IK, Shih WJ, de Cani JS (1990). “Group sequential designs using a family of type I error probability spending functions.” *Statistics in Medicine*, **9**(12), 1439–1445. doi:[10.1002/sim.4780091207](https://doi.org/10.1002/sim.4780091207).
- Jennison C, Turnbull BW (1999). *Group Sequential Methods with Applications to Clinical Trials*. CRC Press.
- Keene ON, Roger JH, Hartley BF, Kenward MG (2014). “Missing data sensitivity analysis for recurrent event data using controlled imputation.” *Pharmaceutical Statistics*, **13**(4), 258–264. doi:[10.1002/pst.1624](https://doi.org/10.1002/pst.1624).
- Lipson DA, Barnhart F, Brealey N, Brooks J, Criner GJ, Day NC, Dransfield MT, Halpin DMG, Han MK, Jones CE, *et al.* (2018). “Once-daily single-inhaler triple versus dual therapy in patients with COPD.” *New England Journal of Medicine*, **378**(18), 1671–1680. doi:[10.1056/NEJMoa1713901](https://doi.org/10.1056/NEJMoa1713901).
- McMurray JJV, Packer M, Desai AS, Gong J, Lefkowitz MP, Rizkala AR, Rouleau JL, Shi VC, Solomon SD, Swedberg K, Zile MR (2014). “Angiotensin–neprilysin inhibition versus enalapril in heart failure.” *New England Journal of Medicine*, **371**(11), 993–1004. doi:[10.1056/NEJMoa1409077](https://doi.org/10.1056/NEJMoa1409077).
- Mütze T, Glimm E, Schmidli H, Friede T (2019). “Group sequential designs for negative binomial outcomes.” *Statistical Methods in Medical Research*, **28**(8), 2326–2347. doi:[10.1177/0962280218773115](https://doi.org/10.1177/0962280218773115).
- National Research Council (2010). *The Prevention and Treatment of Missing Data in Clinical Trials*. National Academies Press, Washington, DC. doi:[10.17226/12955](https://doi.org/10.17226/12955).
- R Core Team (2025). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- Rubin DB (1976). “Inference and Missing Data.” *Biometrika*, **63**(3), 581–592. doi:[10.1093/biomet/63.3.581](https://doi.org/10.1093/biomet/63.3.581).
- Venables WN, Ripley BD (2002). *Modern Applied Statistics with S*. Fourth edition. Springer, New York. ISBN 0-387-95457-0, URL <https://www.stats.ox.ac.uk/pub/MASS4/>.
- Wedzicha JA, Banerji D, Chapman KR, Vestbo J, Roche N, Ayers RT, Thach C, Fogel R, Patalano F, Vogelmeier CF (2016). “Indacaterol–glycopyrronium versus salmeterol–fluticasone for COPD.” *New England Journal of Medicine*, **374**(23), 2222–2234. doi:[10.1056/NEJMoa1516385](https://doi.org/10.1056/NEJMoa1516385).

- Wickham H (2016). *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. ISBN 978-3-319-24277-4. URL <https://ggplot2.tidyverse.org>.
- Wickham H, Hesselberth J, Salmon M, Roy O, Brueggemann S (2025a). *pkgdown: Make Static HTML Documentation for a Package*. doi:10.32614/CRAN.package.pkgdown. R package version 2.2.0, URL <https://CRAN.R-project.org/package=pkgdown>.
- Wickham H, Pedersen TL, Seidel D (2025b). *scales: Scale Functions for Visualization*. doi:10.32614/CRAN.package.scales. R package version 1.4.0, URL <https://CRAN.R-project.org/package=scales>.
- Xie Y (2025). *knitr: A General-Purpose Package for Dynamic Report Generation in R*. R package version 1.51, URL <https://yihui.org/knitr/>.
- Zhu H, Lakkis H (2014). “Sample size calculation for comparing two negative binomial rates.” *Statistics in Medicine*, **33**(3), 376–387. doi:10.1002/sim.5947.

Affiliation:

Keaven Anderson
Merck & Co., Inc., Rahway, NJ, USA
E-mail: keaven_anderson@merck.com

Hongtao Zhang
Merck & Co., Inc., Rahway, NJ, USA

Andrea Maes
Merck & Co., Inc., Rahway, NJ, USA